

Venture capital firms spend a meaningful share of their week on sourcing and triage: scanning inbound, parsing newsletters, searching databases, monitoring founder moves, and trying to catch early signals before a round is obvious. The work is repetitive, tool-fragmented, and highly thesis-dependent. What is a strong lead for one fund is pure noise for another, yet most workflows default to generic alerts, broad keyword searches, and manual spreadsheets.

Problem statement: build a sourcing system that reliably turns the fund's unique thesis into an always-on discovery workflow. The system should reduce noise, surface high-signal companies earlier, and make every recommendation explainable. Today, teams lose time to duplicates, inconsistent tagging, shallow profiles, and the inability to translate "what we invest in" into durable filters and repeatable scouting.

Solution we are thinking about: a "precision AI scout for VCs" that combines (1) a modern intelligence interface for search, filters, profiles, lists, and saved searches, and (2) live enrichment that can pull and summarize public web data on demand. The user experience should feel like a real product: discover companies through fast search + faceted filters, open a profile, enrich it with real website content, generate an explainable score tied to the thesis, then take action (save to list, add notes, share/export). The system is designed to be highly customizable per fund (thesis-first), while remaining transparent about why a company matched and what sources were used.

Competitors / references: Harmonic ([harmonic.ai](https://harmonic.ai)) for workflow inspiration; Cardinal ([trycardinal.ai](https://trycardinal.ai)) as an example of thesis-oriented scouting; broader market references include PitchBook, Crunchbase, CB Insights, Affinity, Dealroom, Tracxn, and SourceScrub. These are reference points for interface patterns and data workflows, not an endorsement or a claim of feature parity.